Dec.’22

MIT Guest Lecture

Qualcomm

Qualcomm
AI research

AI Model Efficiency Toolkit (AIMET)

Chirag Patel
Principal Eng./Mgr.
Qualcomm AI Research

The challenge of AI workloads

- Very compute intensive
- Large, complicated neural network models
- Complex concurrencies
- Real-time
- Always-on

Power and thermal efficiency are essential for on-device AI

Constrained mobile environment

- Must be thermally efficient for sleek, ultra-light designs
- Requires long battery life for all-day use
- Storage/memory bandwidth limitations

Leading machine learning research for edge AI

across the entire spectrum of topics

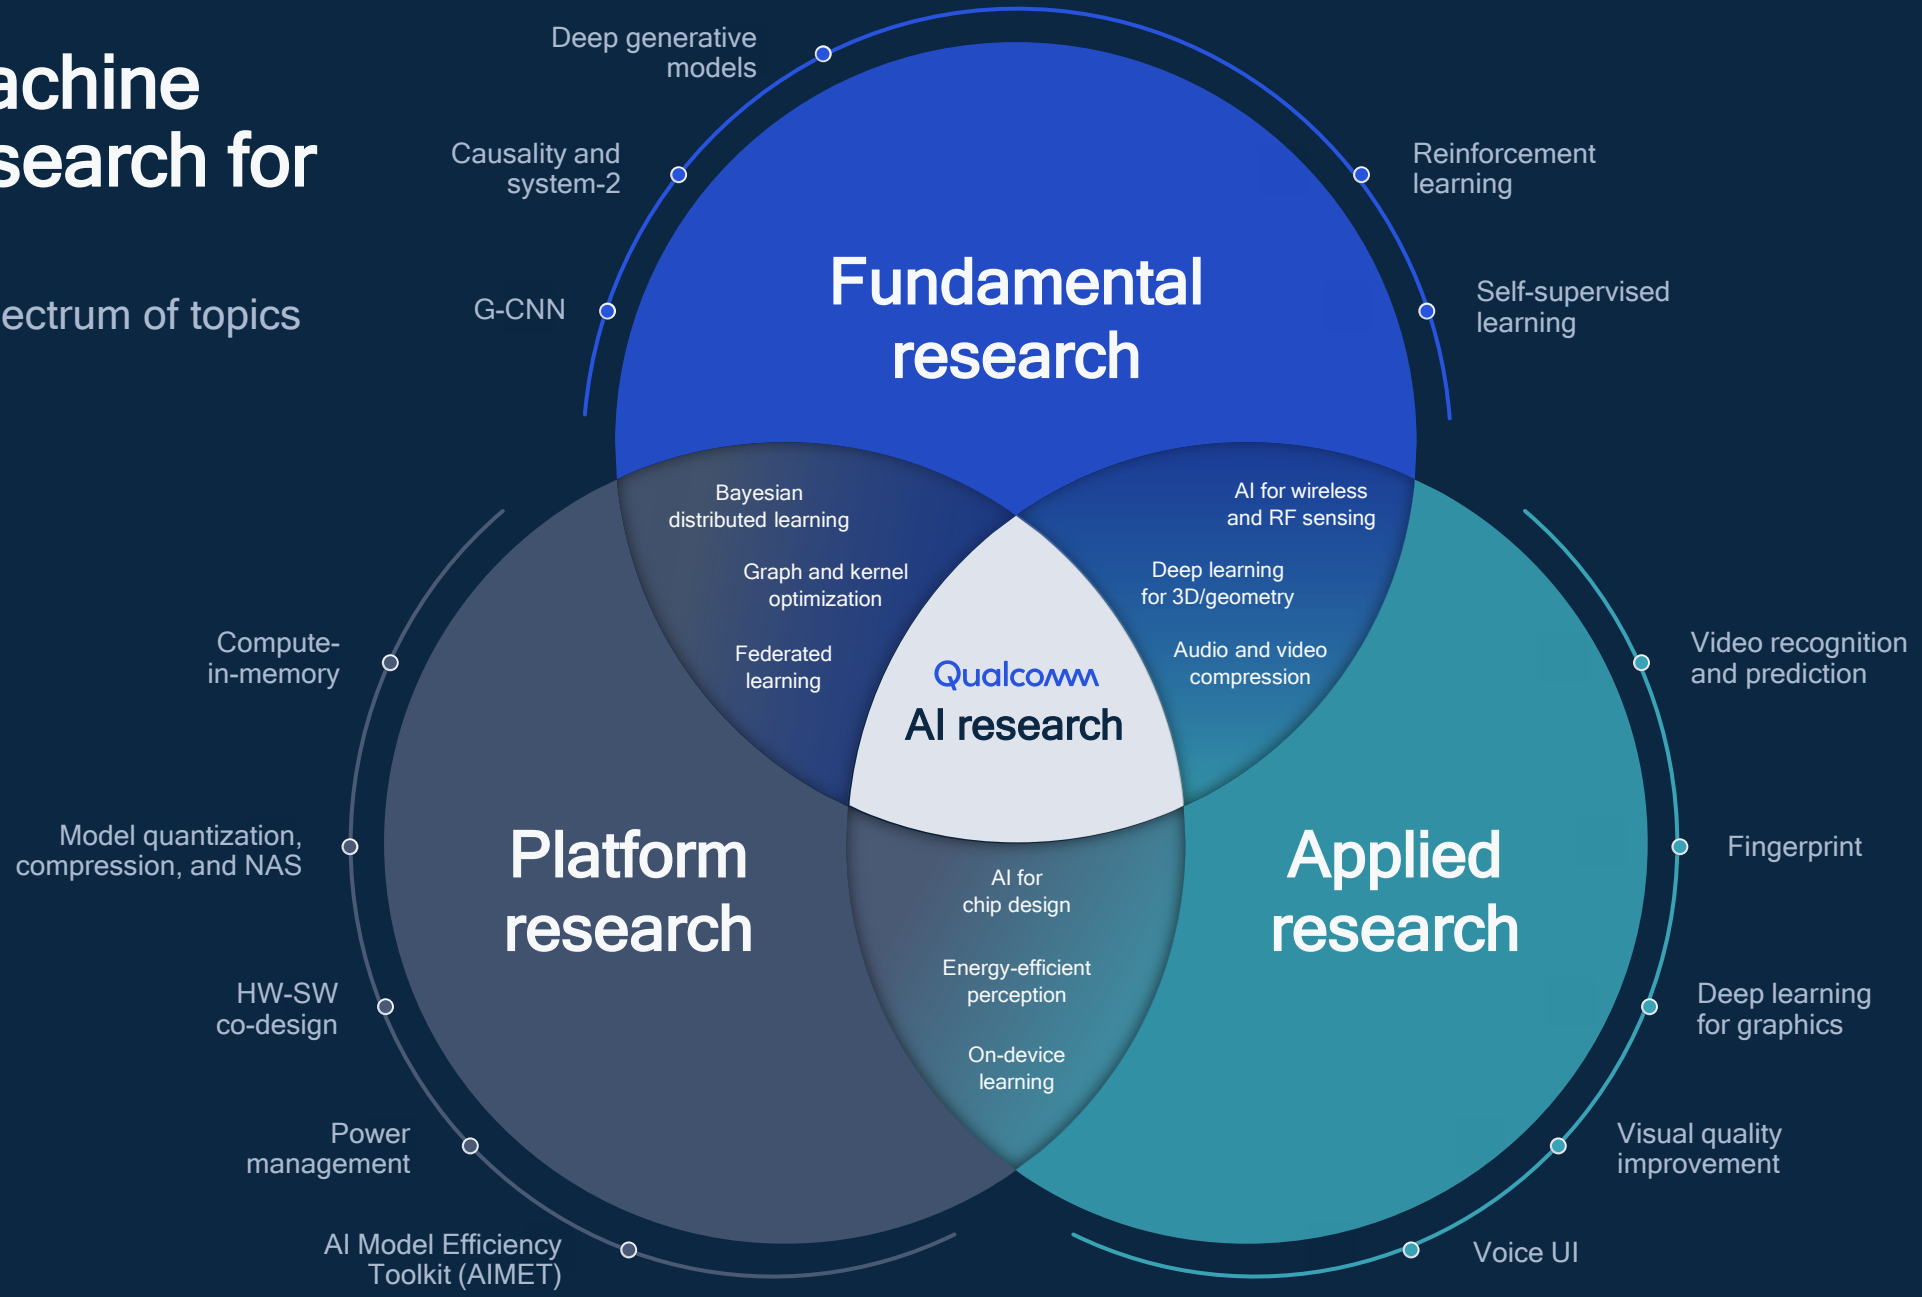

AIMET is a product of Qualcomm Innovation Center, Inc.

Model quantization

Invented the best techniques for fast deployment of 8-bit quantization

Best power-efficiency toolkit in the industry

On-device learning

Invented continuous learning techniques for SOTA on-device voice-UI

First demonstration of 30% improvement to keyword spotting

Federated learning

Invented methods for combining differential privacy and compression

First end-to-end research software framework deployable on mobile

Video semantic segmentation

Top the Cityscape leaderboard with loss function innovation for boundary-awareness

First real-time SS at FHD on mobile

AI Firsts

Brought to you by Qualcomm AI Research

Group equivariant CNN

Pioneer for rotational equivariance; best paper at ICLR'18

First G-CNN segmentation for health on mobile

AI for wireless

Invented neural augmentation to enhance physical layer algorithms

First weakly supervised method for real-world passive RF sensing

Video super resolution

Full stack optimization for visual quality improvement at 4K resolution

First 4K SR at 100+ FPS on mobile

Neural video compression

Invented instance-adaptive for SOTA performance & new deployment scenarios

First real-time HD decoding on mobile

SOTA: State-of-the-art; Cityscapes Benchmark: https://www.cityscapes-dataset.com/

Holistic model efficiency research

Multiple axes to shrink AI models and efficiently run them on hardware

Quantization

Learning to reduce bit-precision while keeping desired accuracy

Compilation

Learning to compile AI models for efficient hardware execution

Conditional compute

Learning to execute only parts of a large inference model based on the input

Neural architecture search

Learning to design smaller neural networks that are on par or outperform hand-designed architectures on real hardware

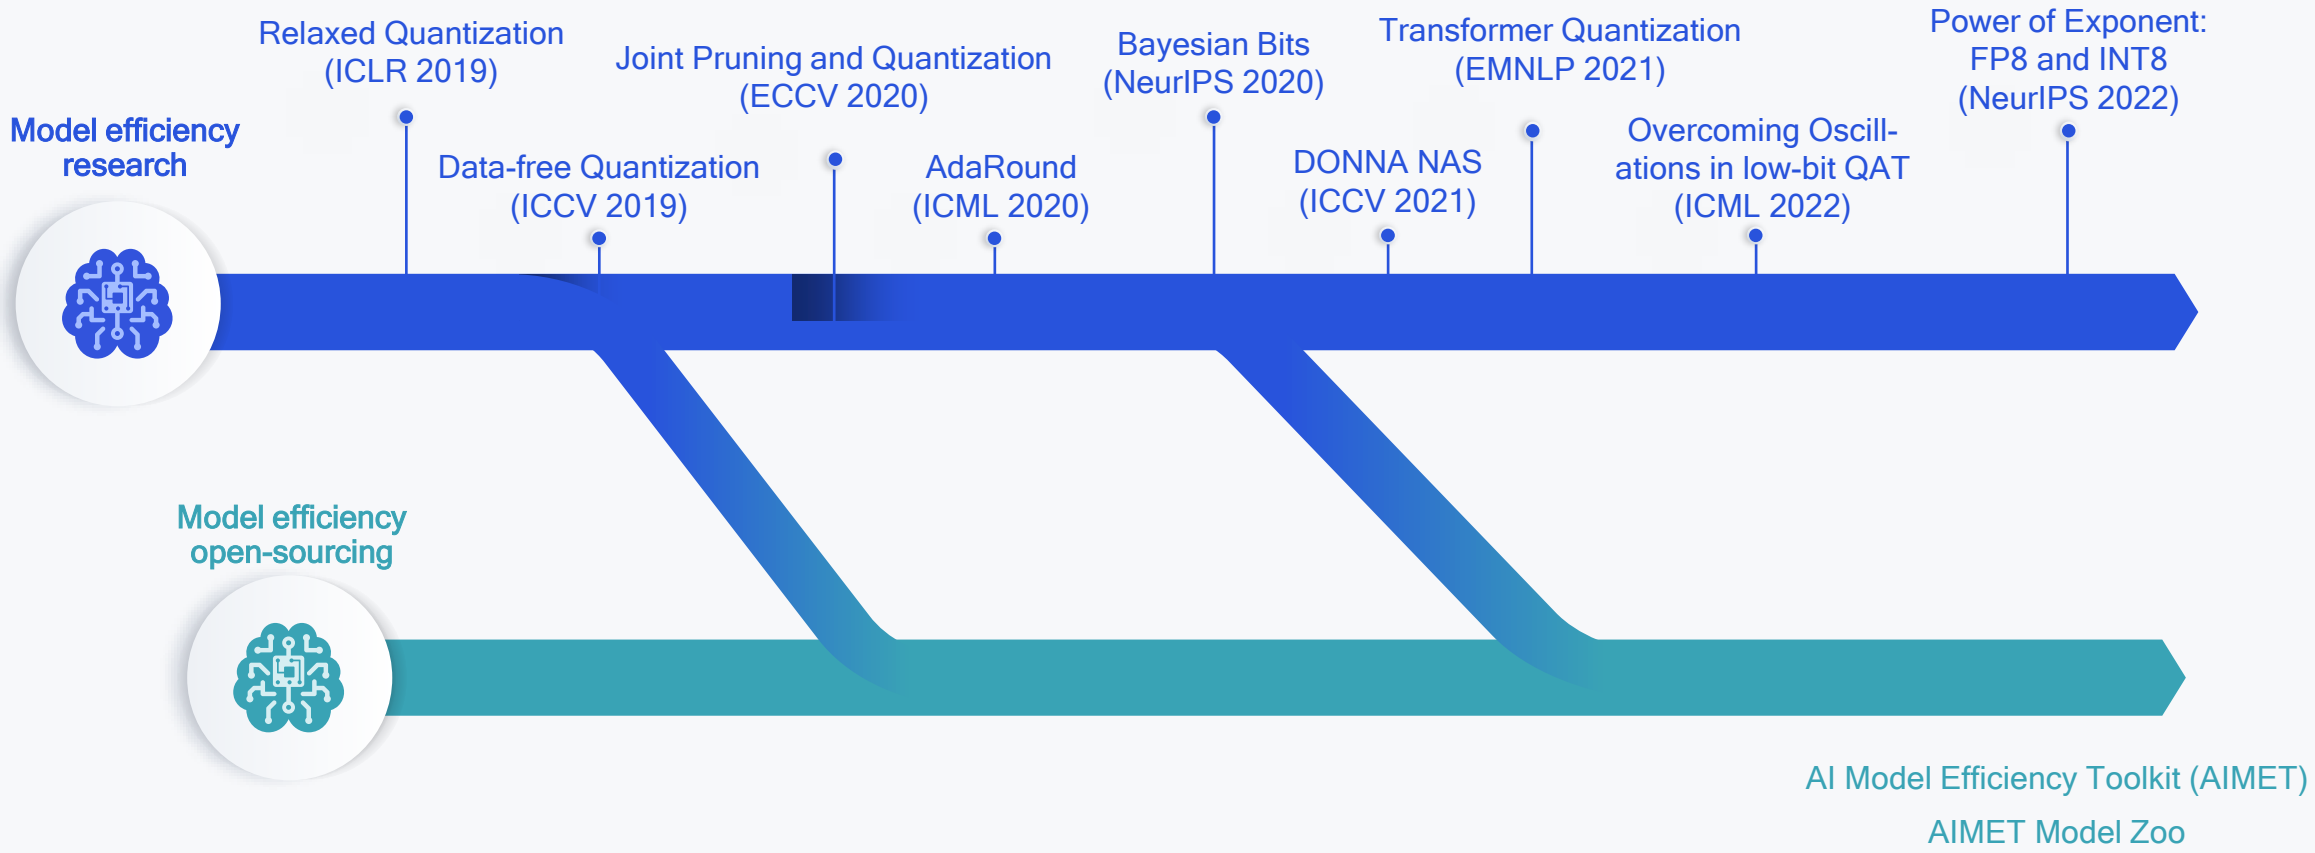

Driving the industry toward integer inference and power-efficient AI

Leading model efficiency research and fast commercialization

AIMET makes AI models small

State-of-the-art quantization and compression techniques from Qualcomm AI Research

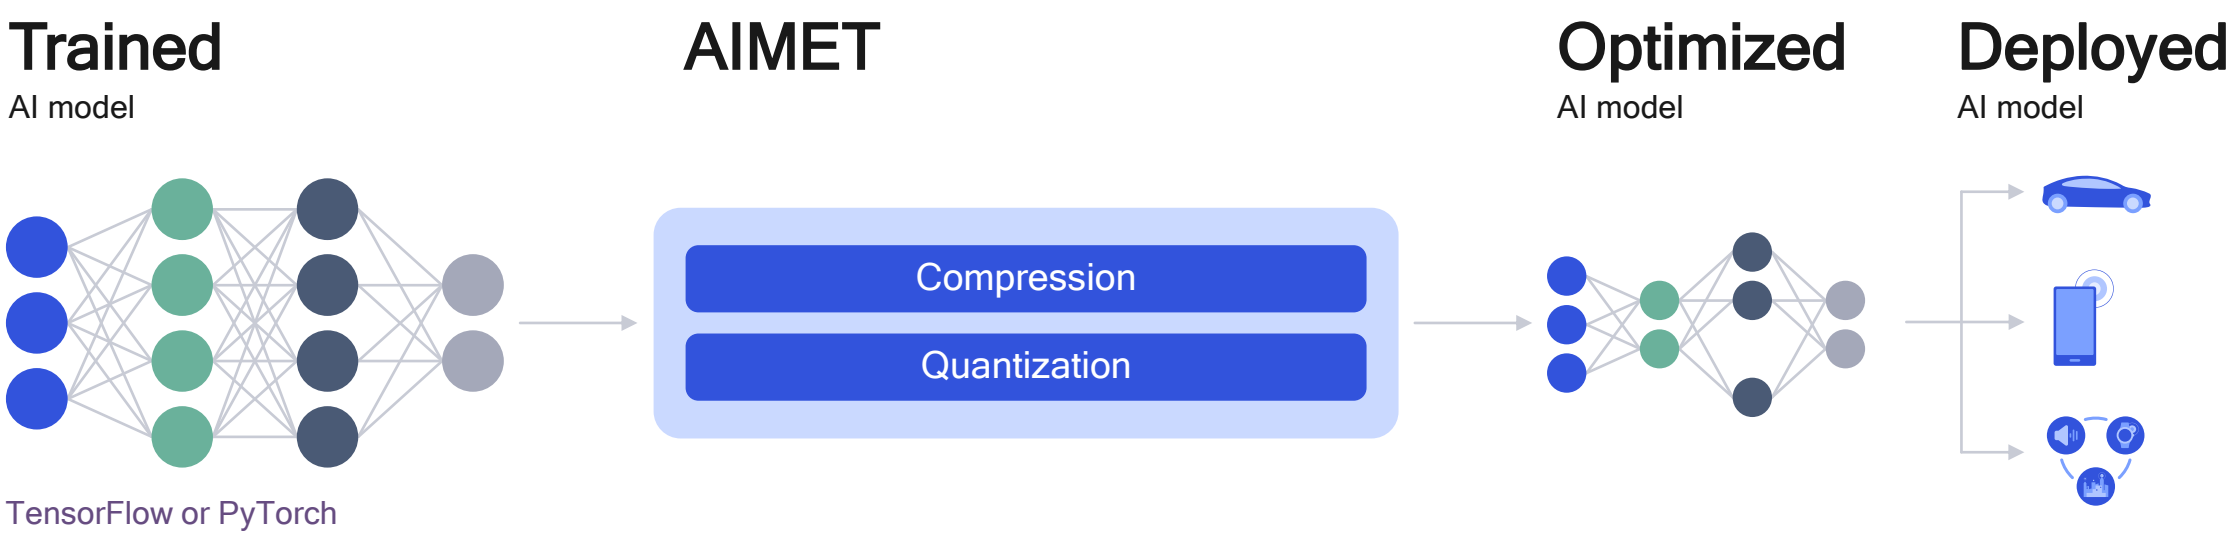

Github: https://github.com/quic/aimet

Features:

- State-of-the-art network compression tools
- State-of-the-art quantization tools
- Support for both TensorFlow and PyTorch
- Benchmarks and tests for many models
- Developed by professional software developers

AIMET

Providing advanced model efficiency features and benefits

Benefits

Lower power

Lower memory bandwidth

Maintains model accuracy

Lower storage

Higher performance

Simple ease of use

Features

Quantization

State-of-the-art INT8 and INT4 performance

Quantization simulation

Post-training quantization (PTQ) methods:

- Data-Free Quantization
- Adaptive Rounding (AdaRound),
- Automatic Mixed Precision (AMP)
- AutoQuant

Quantization-aware training (QAT)

Compression

Efficient tensor decomposition and removal of redundant channels in convolution layers

Spatial singular value decomposition (SVD)

Channel pruning

Visualization

Analysis tools for drawing insights for quantization and compression

Weight ranges

Per-layer compression sensitivity

AdaRound: Adaptive Rounding for Better Quantization

ICML'20 paper

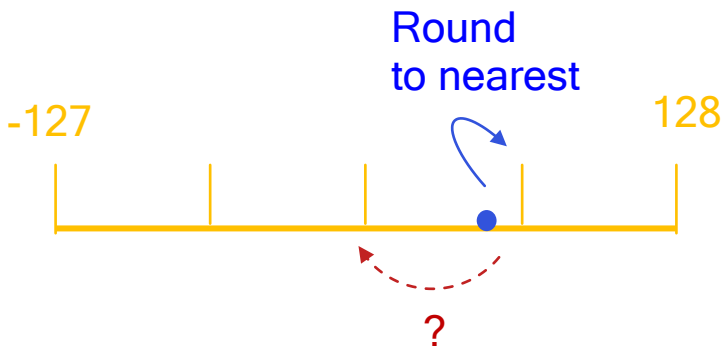

Rounding-to-the-nearest is not optimal

Object Detection

Configuration	mAP
Floating point	82.20
Nearest Rounding - 8-bit weights, 8-bit activations	49.85
AdaRound - 8-bit weights, 8-bit activations	81.21

mAP: Mean Average Precision

Semantic Seg. (Deeplabv3)

Configuration	mIOU
Floating point	72.94
Nearest Rounding - 4-bit weights, 8-bit activations	6.09
AdaRound - 4-bit weights, 8-bit activations	70.86

mIOU: Mean Intersection Over Union

AdaRound optimizes the network weights without model fine-tuning

$$\arg\min_{\mathbf{V}} \left\| \mathbf{W}\mathbf{x} - \widetilde{\mathbf{W}}\mathbf{x} \right\|_F^2 + \lambda f_{reg}(\mathbf{V})$$

AdaRound Results

- Poor baseline INT8 quantization performance
- AdaRound performance within 1% of FP32

INT8, Baseline

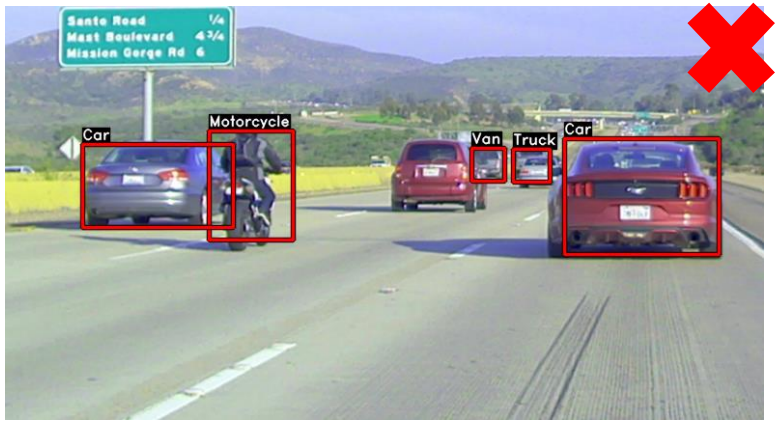

INT8, AdaRound

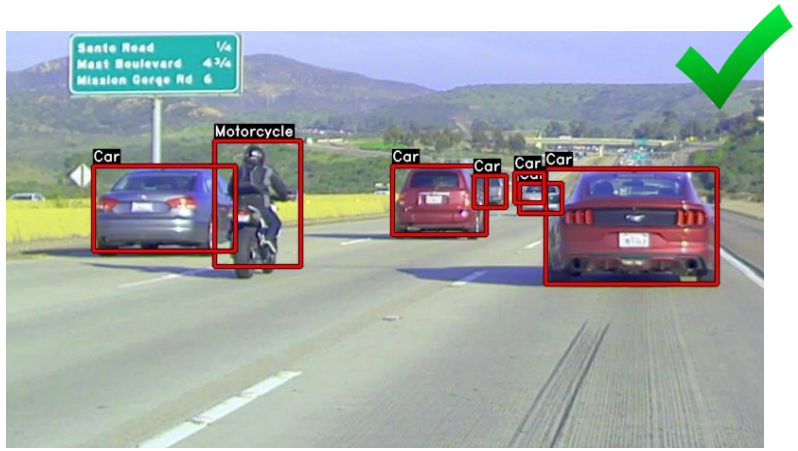

AutoQuant simplifies post-training quantization

- Analyzes the model
- Applies the best sequence of already existing post-training quantization (PTQ) features
- Returns the best accuracy model with analysis report
- A simple, blackbox, push-button solution

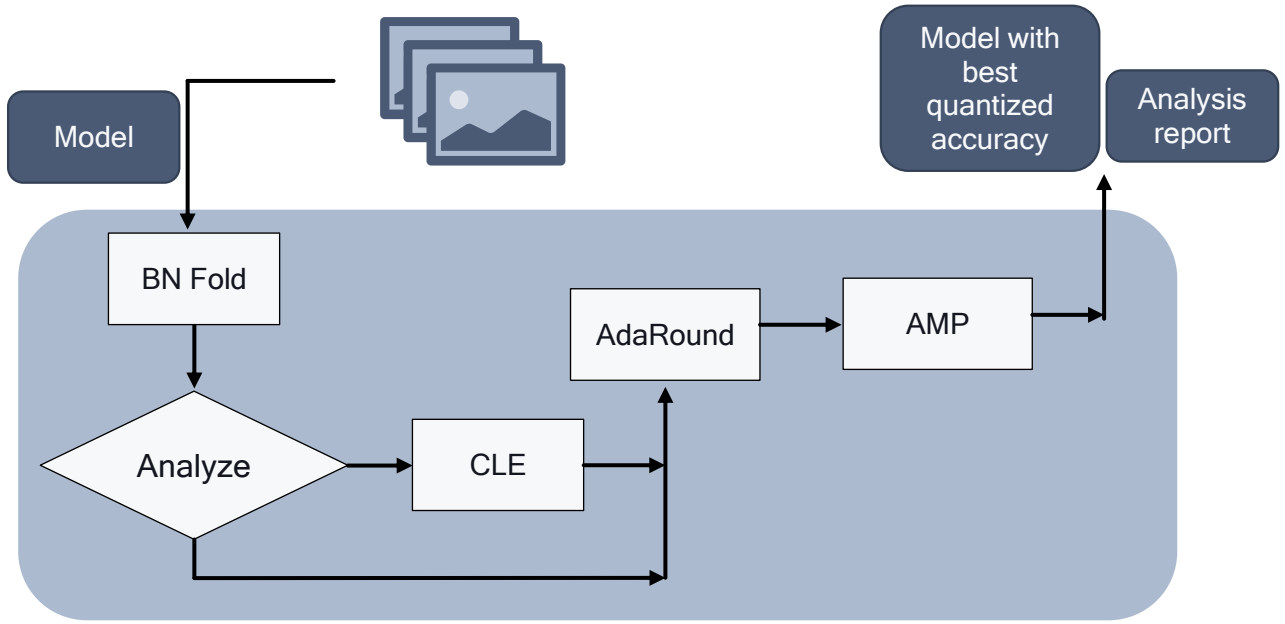

AIMET Quantization Aware Training

Simulate quantization noise in the forward pass and fine-tune for improved robustness

Provides accurate prediction of on-target performance by HW/run-time awareness

INT8 performance typically within 0.5-1% of FP32 performance

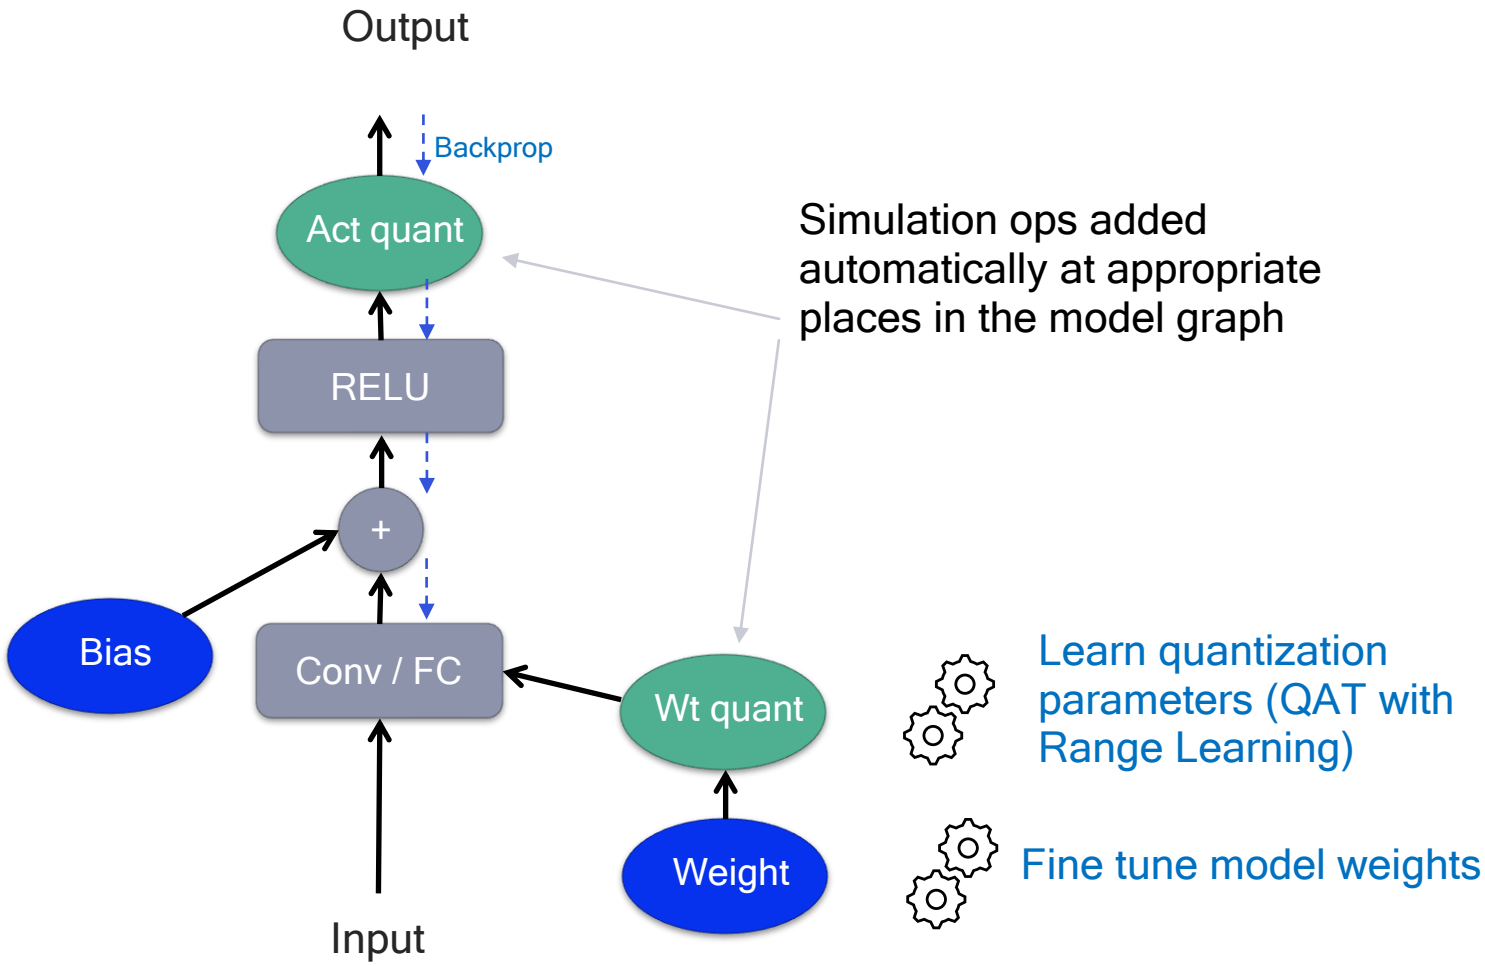

TensorFlow

<1%
Loss in accuracy*

PyTorch

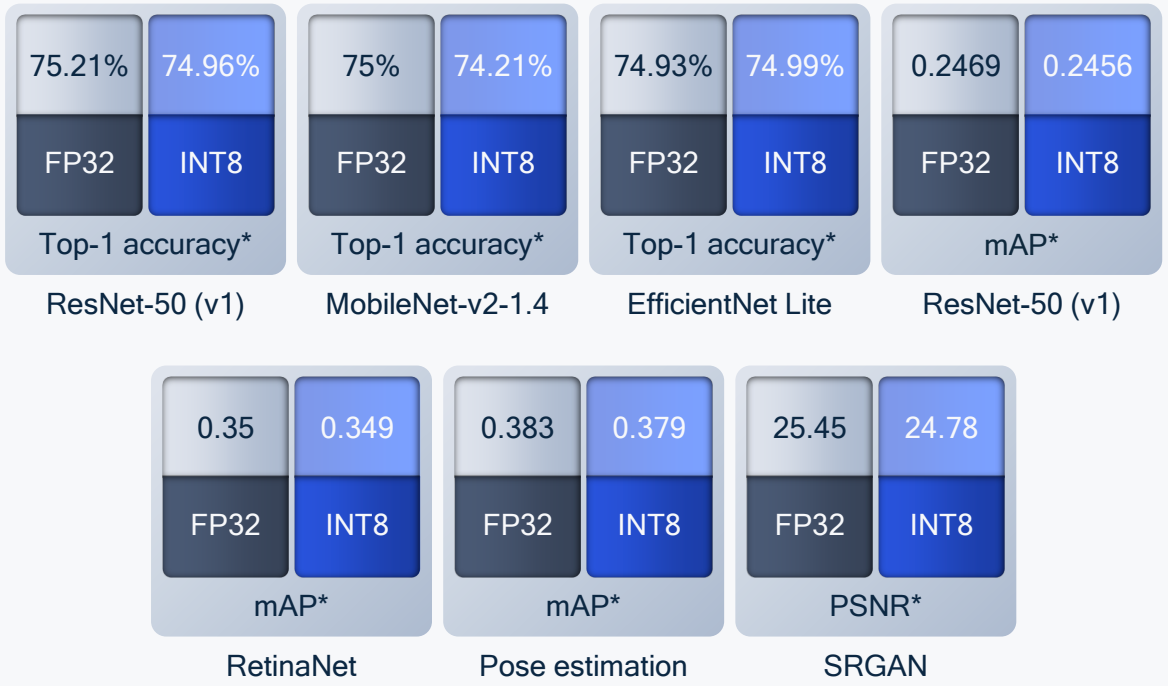

AIMET Model Zoo includes popular quantized AI models

Accuracy is maintained for INT8 and INT4 models – less than 1% loss*

*: Comparison between FP32 model and INT8 model quantized with AIMET. For further details, check out: https://github.com/quic/aimet-model-zoo/

Transformer Quantization

Model	FP32	INT8
BERT-base-uncased	82.73 (GLUE)	82.53
DistilBERT-base-uncased	80.35 (GLUE)	79.81
mobileBERT	81.24 (GLUE)	81.27
VIT (vision transformer)	81.30	81.50

AIMET quantizes transformers with high accuracy, comparable to FP32

Comparison between FP32 model and model quantized with AIMET

AIMET enables accurate INT4 (4-bit weights, 8-bit activations) for wide range of use cases

Task	Model	FP32	INT W4A8
Classification	ResNet50	76.10%	75.4%
	ResNet18	69.75%	68.96%
	EfficientNet-Lite	75.31%	74.33%
	Regnext	78.30%	77.20%
Segmentation	Deeplabv3 (RN-50)	76.07%	75.91%
Super-resolution	ABPN	31.97 dB	31.67 (dB)
Pose detection	PoseNet (HRNet-32)	0.765	0.763

With better PTQ and QAT techniques, more models will achieve better power efficiency

AIMET

State-of-the-art quantization and compression techniques

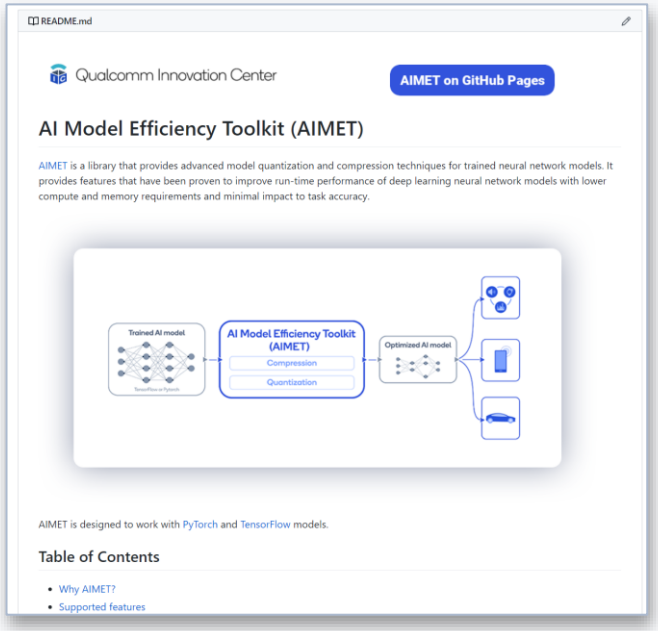

github.com/quic/aimet

AIMET Model Zoo

Accurate pre-trained 8-bit quantized models

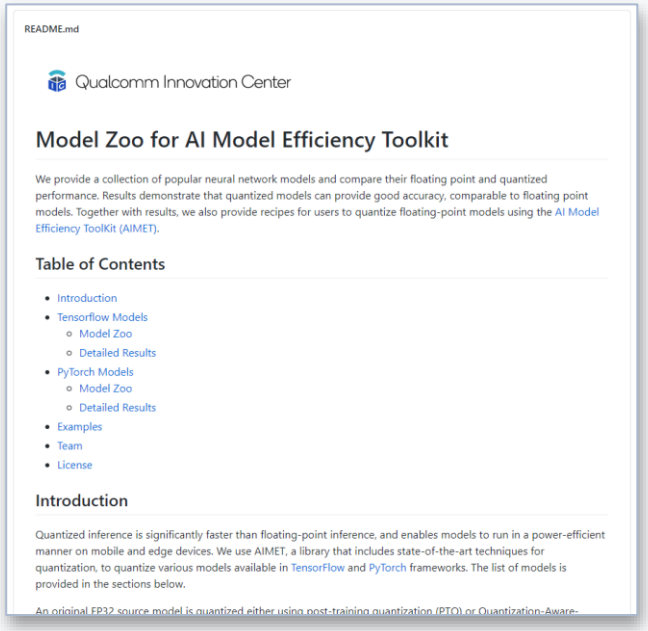

github.com/quic/aimet-model-zoo

Quantization whitepaper

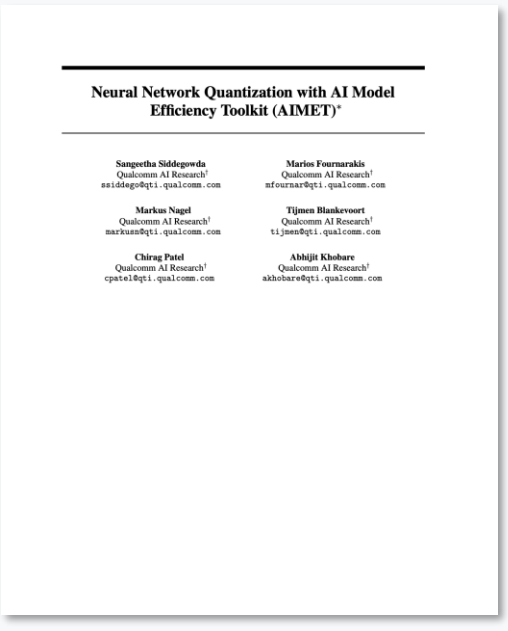

arxiv.org/abs/2201.08442

Explore our open-source projects and tools

Qualcomm AI Stack

Tools:

Qualcomm AI Model Studio

- AIMET
- AIMET Model Zoo
- NAS
- Model analyzers

Infrastructure:

Prometheus

kubernetes

docker

AI Frameworks, Runtimes

Developer Libraries and Services

System Software

OS

AI Frameworks

TensorFlow PyTorch ONNX

AI Runtimes

Qualcomm® Neural Processing SDK | ONNX RUNTIME | TF Lite Micro | Direct ML | TF Lite

Qualcomm® AI Engine Direct (QNN)

Math Libraries | Compilers | Virtual platforms

Profilers & Debuggers | Programming Languages | Core Libraries

System Interface | SoC, accelerator drivers | Emulation Support

android | Zephyr™ | ubuntu® | CentOS | QNX

Platforms

Smartphones | XR | ACPC | IoT | Robotics | Auto | Cloud







Snapdragon® 8 Gen 2 Mobile Platform Qualcomm® AI Engine

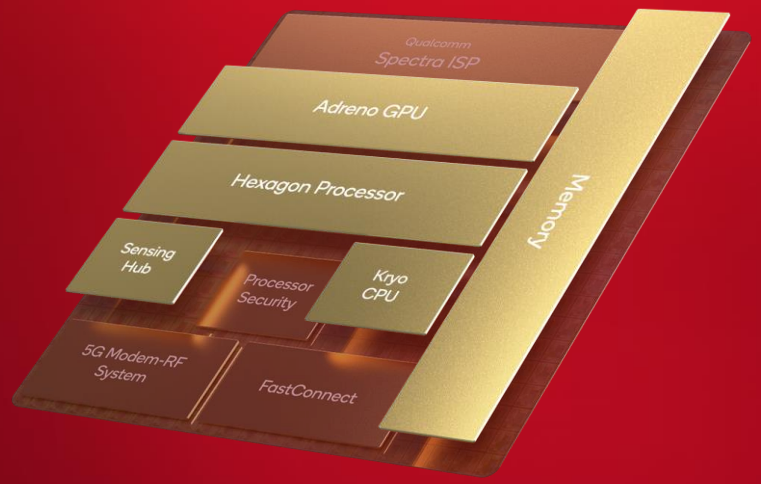

More Hardware Acceleration

2X Tensor Accelerator Performance

up to **4.35X** Performance improvement*

60% More power efficient

Multi-language Translation

World's 1ST **INT4** support

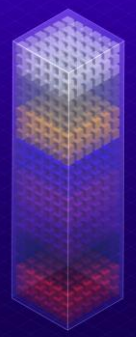

ArcSoft® AI Cinematic Mode

Qualcomm AI Stack

UNREAL ENGINE

Feature updates
Performance Improvements

AI bot plug-in

INT4 Support

Qualcomm Sensing Hub

Dual AI Processor

Always-sensing Camera

Thank you

Qualcomm

Follow us on:

For more information, visit us at:

qualcomm.com & qualcomm.com/blog